

# Additive Attention Mechanism

---

## Overview

Additive Attention, introduced by Bahdanau et al. (2015), is a mechanism commonly used in sequence-to-sequence models. It computes attention scores between a query and a set of keys, then uses these scores to weight the corresponding values. The core of the mechanism is an energy function that computes compatibility between the query and keys.

---

## Why Use Additive Attention?

- **Alignment Learning:** Helps in determining which parts of the input sequence are important for producing the output.
  - **Sequence-to-Sequence Models:** Typically used in tasks such as machine translation, speech recognition, and text summarization.
  - **Improved Contextualization:** Adds the ability to focus on different parts of the input sequence for different outputs.
- 

## Prerequisites

- **Python 3.x:** Ensure Python 3.x is installed.
- **PyTorch:** Install PyTorch compatible with your system.

```
pip install torch
```

---

## Code Description

### 1. Data Preparation:

Here, random data for input sequences, a single query vector, and keys/values for the attention mechanism are generated.

```
X = np.random.randn(10, 20) # Random data for input (10 sequences, 20 features)
query = torch.randn(1, 20) # A single query vector
key = torch.randn(10, 20) # Key for all sequences
value = torch.randn(10, 20) # Value for all sequences
```

- `X` : The input data (for example, the output from an encoder).
- `query` : A vector representing the query (typically the output from a decoder).
- `key` : The set of keys, representing all possible input sequences.
- `value` : The corresponding values for each key.

### 2. Additive Attention Class:

The `AdditiveAttention` class defines the layers and forward pass of the attention mechanism.

```
class AdditiveAttention(nn.Module):
    def __init__(self, input_dim, hidden_dim):
        super(AdditiveAttention, self).__init__()
        self.query_layer = nn.Linear(input_dim, hidden_dim)
        self.key_layer = nn.Linear(input_dim, hidden_dim)
        self.energy_layer = nn.Linear(hidden_dim, 1)

    def forward(self, query, key, value):
        query = self.query_layer(query).unsqueeze(0)
        key = self.key_layer(key)
        energy = torch.tanh(query + key)
        attention_weights = torch.softmax(self.energy_layer(energy).squeeze(), dim=-1)
        weighted_sum = torch.sum(attention_weights.unsqueeze(1) * value, dim=-1)
        return weighted_sum, attention_weights
```

- `query_layer` and `key_layer` : Linear layers to transform the query and key vectors into a shared hidden space.
- `energy_layer` : Computes a scalar energy score from the sum of transformed query and key vectors.
- `forward` : The attention mechanism that computes the attention weights and the weighted sum of values.

### 3. Attention Computation:

The attention mechanism is applied to the input data, computing the weighted sum and attention weights.

```
attention = AdditiveAttention(input_dim=20, hidden_dim=32)
weighted_sum, attention_weights = attention(query, key, value)
```

- The `weighted_sum` is the final output, which is a weighted sum of the value vectors based on the computed attention weights.
- The `attention_weights` represent how much focus each input sequence (key-value pair) receives for this query.

### 4. Output:

The weighted sum and attention weights are printed for inspection.

```
print("Weighted Sum:", weighted_sum)
print("Attention Weights:", attention_weights)
```

---

## Expected Outputs

- **Weighted Sum:** The resulting weighted sum of the value vectors based on attention scores.
  - **Attention Weights:** A tensor indicating how much attention is given to each key-value pair in the sequence.
- 

## Use Cases

- **Machine Translation:** Additive attention can be applied in machine translation models to focus on relevant parts of the input sentence when generating the translated output.
  - **Speech Recognition:** In speech-to-text models, attention mechanisms help focus on specific parts of the speech signal when transcribing it to text.
  - **Text Summarization:** For generating summaries, the model can focus on the most relevant portions of the document.
- 

## Future Enhancements

1. **Multi-Head Attention:** Implement multi-head attention to learn different attention representations from different parts of the input.
  2. **Scalability:** Improve the mechanism to handle longer sequences efficiently (e.g., by using sparse attention or performing attention in chunks).
  3. **Visualization:** Add functionality to visualize the attention weights to understand which parts of the input are being focused on during prediction.
- 

## References

- Bahdanau, D., Cho, K., & Bengio, Y. (2015). Neural Machine Translation by Jointly Learning to Align and Translate. [arXiv:1409.0473](https://arxiv.org/abs/1409.0473).