

# Denoising Autoencoder (DAE)

---

## Overview

A **Denoising Autoencoder (DAE)** is a type of autoencoder that learns to reconstruct the original, clean input from a corrupted version of the data. This approach can help the model learn more robust and generalized representations by forcing it to focus on the essential features of the data, ignoring noise and irrelevant information.

The Denoising Autoencoder is a very effective tool for unsupervised learning, denoising applications, and improving generalization capabilities in neural networks.

---

## Why Use Denoising Autoencoders?

Denoising Autoencoders are beneficial for the following reasons:

- **Noise Robustness:** DAE forces the network to learn a more robust representation of the data by introducing noise and requiring the network to recover the clean version of the data.
  - **Pre-training for Deep Networks:** DAEs can serve as a pre-training mechanism for deep neural networks, leading to better performance when fine-tuned for specific tasks.
  - **Unsupervised Learning:** Denoising Autoencoders are effective for learning representations without needing labeled data, making them useful for unsupervised learning tasks.
- 

## Prerequisites

- **Python 3.x:** Ensure Python 3.x is installed.
- **PyTorch:** Install PyTorch compatible with your system.
- **scikit-learn:** For dataset generation and preprocessing.

```
pip install scikit-learn
```

- **matplotlib:** For plotting results (optional).

```
pip install matplotlib
```

---

## Files Included

- **denoising\_autoencoder.py** : Contains the implementation of the Denoising Autoencoder model.
  - **train.py** : Script to train the Denoising Autoencoder model on the synthetic dataset.
  - **plot\_results.py** : Script to visualize the training results (optional).
- 

## Code Description

## 1. Data Generation and Preprocessing:

```
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
import numpy as np

X, y = make_classification(n_samples=1000, n_features=20, n_classes=2, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

mean = np.mean(X_train, axis=0)
std = np.std(X_train, axis=0)
X_train = (X_train - mean) / std
X_test = (X_test - mean) / std
```

This part generates a synthetic dataset with 1,000 samples and 20 features, then standardizes the dataset by subtracting the mean and dividing by the standard deviation.

## 2. Denoising Autoencoder Model Definition:

```
import torch
import torch.nn as nn
import torch.optim as optim

class DenoisingAutoencoder(nn.Module):
    def __init__(self, input_dim, hidden_dim):
        super(DenoisingAutoencoder, self).__init__()
        self.encoder = nn.Linear(input_dim, hidden_dim)
        self.decoder = nn.Linear(hidden_dim, input_dim)

    def forward(self, x):
        encoded = torch.relu(self.encoder(x))
        decoded = torch.sigmoid(self.decoder(encoded))
        return decoded
```

The `DenoisingAutoencoder` class consists of an encoder and a decoder. The encoder compresses the input data into a lower-dimensional space, while the decoder reconstructs the original data from this compressed representation.

## 3. Noise Addition:

```
def add_noise(x, noise_factor=0.2):
    noise = torch.randn_like(x) * noise_factor
    return torch.clamp(x + noise, 0, 1)
```

The `add_noise` function adds Gaussian noise to the input data, simulating a corrupted version of the original input. The `torch.clamp` ensures that the noisy values stay within the range [0, 1].

## 4. Model Training:

```
model = DenoisingAutoencoder(input_dim=X_train.shape[1], hidden_dim=8)
optimizer = optim.Adam(model.parameters(), lr=0.001)
loss_fn = nn.MSELoss()

for epoch in range(50):
```

```
optimizer.zero_grad()
noisy_input = add_noise(torch.tensor(X_train, dtype=torch.float32))
y_pred = model(noisy_input)
loss = loss_fn(y_pred, torch.tensor(X_train, dtype=torch.float32))
loss.backward()
optimizer.step()

if (epoch + 1) % 10 == 0:
    print(f"Epoch [{epoch+1}/50], Loss: {loss.item():.4f}")
```

The model is trained using the Adam optimizer with a Mean Squared Error (MSE) loss. During each epoch, noise is added to the input, and the model tries to reconstruct the original, clean input. The loss is calculated as the difference between the model's output and the true input.

## 5. Model Evaluation (optional):

```
from sklearn.metrics import mean_squared_error

y_pred = model(torch.tensor(X_test, dtype=torch.float32))
mse = mean_squared_error(X_test, y_pred.detach().numpy())
print(f"Test MSE: {mse:.4f}")
```

After training, the model is evaluated using Mean Squared Error (MSE) on the test set. This provides insight into the model's ability to reconstruct the noisy data.

## 6. Visualization (optional):

```
import matplotlib.pyplot as plt

# Plot reconstructed data vs. original data
plt.plot(X_test[0], label='Original')
plt.plot(y_pred[0].detach().numpy(), label='Reconstructed')
plt.legend()
plt.title('Original vs Reconstructed Data')
plt.show()
```

This block visualizes a comparison between the original and reconstructed data. It's helpful to see how well the model denoised the input.

---

# Expected Outputs

- **Training Progress:** The console will display the loss at every 10th epoch, reflecting the model's convergence over time.
  - **Test Performance:** After training, the MSE on the test set will be printed, showing the model's reconstruction accuracy.
  - **Visualization:** A plot comparing the original and reconstructed data for a sample from the test set will be shown.
- 

# Use Cases

- **Data Denoising:** Removing noise from images, signals, or other data types by training the model on noisy inputs.
  - **Feature Learning:** Learning useful and noise-robust representations from data in an unsupervised way.
  - **Pre-training:** Used as a pre-training technique for initializing weights in deep neural networks, which can then be fine-tuned on specific tasks.
-

## Future Enhancements

1. **Hyperparameter Optimization:** Experiment with various architectures, noise types, and noise levels to improve model performance.
  2. **Advanced Denoising Techniques:** Integrate other types of noise or explore different reconstruction loss functions such as L1 loss or perceptual loss.
  3. **Real-World Applications:** Test the model on real-world datasets, such as noisy image datasets (e.g., MNIST, CIFAR-10).
  4. **Visualize Latent Space:** Use dimensionality reduction techniques like PCA or t-SNE to visualize the learned representations and inspect their separability.
  5. **Integration with Convolutional Networks:** Implement Denoising Autoencoders using convolutional layers for better performance on image data.
- 

## References

- Vincent, P., et al. (2008). Extracting and Composing Robust Features with Denoising Autoencoders. *Proceedings of the 25th International Conference on Machine Learning (ICML)*. [Link](#).
- Bengio, Y., et al. (2013). Learning Deep Architectures for AI. *Foundations and Trends® in Machine Learning*. [Link](#).